

HANDS-ON WITH CLAUDE CODE · 31 JULY 2026 · CIC BERLIN

Build Your Weekly Briefing Agent

A real AI workflow agent that reads your week and hands it back, sorted. **No code.**
Two hours. Yours to keep.

Tobias Nentwig · founder, nyxory · building with Claude Code daily

THE NEXT TWO HOURS

By the end of tonight, **you'll get**

- 01 **What an AI agent actually is**, and why it's not just ChatGPT.
- 02 **How an agent reaches your tools**, your email, your calendar, your notes.
- 03 **How to build one yourself**, tonight, on your own laptop, no code.
- 04 **How to reuse the same idea** for dozens of other workflows.

No jargon yet. We'll take it one step at a time, **everyone leaves knowing where we went.**

WHAT YOU LEAVE WITH

In two hours you don't learn to code. You learn to **delegate to a system you trust.**

You'll walk out with your own agent, running on your laptop, doing a real job — and the confidence to build the next one without us.



SO NO WORD TONIGHT IS A SURPRISE

Six words you'll **know by the end**

Prompt one conversation, you ask, it answers once.

Prompt engineering the craft of the ask — wording, examples, what good looks like.

MCP the connector that plugs a tool — Gmail, Calendar, Notion — into the AI.

Agent repeated work, set up once, runs every time.

Context engineering the craft of what it knows *before* you ask — the files around it. Tonight's real skill.

Claude Code where we build tonight — plain English in, working agent out.

Every time one of these comes up later, you'll already have it. **No catching up.**

WHERE YOU ARE, AND WHERE WE'RE GOING

You've used AI. Tonight you level up.

Chat

You ask, it answers. Text comes back — nothing else happens.



Cowork agent

Hands and a memory. You work on things together, and it gets better at your work every week.

TONIGHT WE'RE HERE



Remote agent

Takes one of those jobs and does it alone, while you're somewhere else.

LATER TONIGHT

A **multi-agent system** — several agents splitting the work between them — sounds like the top of the ladder. It isn't: **that can happen inside any of these three.**

HOW A REMOTE AGENT ACTUALLY GETS BUILT

You don't start one. You **cut one out**.

Your cowork agent

knows you, your customers, your tone, a year of notes

briefing

follow-ups

customer one-pagers

research

CRM notes

outreach



lift one job out

One remote agent

customer one-pagers

Runs on its own.


Share it with the team.

Building a production agent from a blank page is hard. Lifting one job out of an agent that **already does it well with you** is not — the hard part, knowing your work, is already done.

THE DIFFERENCE THAT MATTERS

A prompt **answers once.** An agent **runs every time.**

You don't re-explain yourself each morning. You set it up once — who it works for, what matters, what it can touch — and it does the job again next Monday, and the one after.



WHAT DOES AN AGENT NEED TO DO THE JOB?

Six things. That's all.



Brain

Claude, it reads and decides.



Identity

who it is and what its job is.



Context

who you are and how you sound.



Memory

what it keeps from last time.



Tools

Gmail, Calendar, Notion, what it can reach.



Rubric

your standard for what matters.

You'll set five of these tonight, in **plain English**, not code. **Memory** is the one that grows by itself.

THE ONE THING THAT MAKES IT USEFUL

The difference is **tools**.

Without tools

AI can answer questions, write text, brainstorm ideas.

Helpful, but it can't touch **your** world.

With tools

It can read your Gmail, check your Calendar, create documents, update Notion, search your sheets.

Now it works on **your actual week**.

Which raises the obvious question: **how does it connect to Gmail?** → that's next.

NOT THE SAME AS KNOWING WHO YOU ARE

Memory is what it **keeps**

Context

Who you are, what you do, how you sound. Written once, changes rarely.

"I'm a PM. Short and direct, never bullet points."

Memory

What it holds on to **so the next run is better than the last**. It grows every week.

"Call with Marta, 24 Jul: pilot agreed, SSO still open."

"That draft was too long — she reads on her phone."

"Priya's legal team only answers in writing."

Memory is just **notes in a folder** your agent can read. If you keep notes in Obsidian, that folder already *is* your agent's memory — it's markdown files either way.

WHY EVERY AGENT GETS ITS OWN FOLDER

The folder **is** the memory

Claude Code

Your agent **works** here. It reads the folder before every run, and writes back what it learns.



 **weekly-briefing-agent/**

IDENTITY.md	who it is
me.md	who you are
rubric.md	your standard
vault/	what it learns, run by run



Obsidian

Same folder, **second window**. Read your agent's memory like a wiki — and edit it. Your edit is in its head next run.

Claude itself forgets everything when the chat ends. Whatever should survive must be **a file in the folder** — open the folder, and the agent wakes up knowing it.

Look at it

Type `/memory` in Claude Code — its notebook for the folder you're in. As you work together it adds its own notes there, under `~/.claude/projects/<this folder>/memory/`.

Memory is **folder-specific** — that's why it's one folder per agent. Open the folder → that agent, that memory. Two agents in one folder share one brain, and confuse each other.

IF YOU REMEMBER ONE THING TONIGHT

Every agent is the **same three moves**


READ a source	JUDGE it against a rubric you wrote	WRITE where you already look
-------------------------------------	---	--

Screeener, radar, invoice-sorter, your briefing — all the same skeleton. **Learn it once, you can build any of them.**

HOW AN AGENT TOUCHES THE REAL WORLD

An **API** is a socket. An **MCP** is the plug that fits it.

Your email, your calendar, your notes each expose a socket. An **MCP** is a ready-made plug your agent uses to reach in — read a thread, add a draft — **without you writing any code**. One click connects one plug.



Tonight we connect one live, so you feel it: **an agent reaching into a real inbox.**

SO THE WORD STOPS BEING ABSTRACT

This is what an **API** actually looks like

openholidaysapi.org • free, no key

```
{  
  "name": "Tag der Deutschen Einheit",  
  "startDate": "2026-10-03",  
  "nationwide": true  
}
```

Public holidays in Germany. Written for machines, not for you.

You never have to read this. You tell Claude Code what you want — *"add the next public holiday to my briefing"* — and it deals with the format.

Some tools need a **key** first: a password for your agent. You get it on the tool's own site, usually under **DeveLopers** or **API**, and paste it in once.



Don't believe me? Scan it — that's the **same call, live**, on your phone.

An **MCP** is just this, pre-wired — someone already did the plumbing so you only have to click connect.

DEMO — ONE PLUG, ONE INBOX

I connect my Gmail and let the agent read my week.

One click in the browser. No project to set up, no code. Watch it sort what needs me from what doesn't — in real time, on my real mail.

YOUR TURN — THE WEEKLY BRIEFING AGENT

Reads your week → writes the **replies that matter**

Your week

emails + calendar
(sample data to start)



The agent

judges it against
your rubric




Ready to send

drafted replies · decisions ·
week ahead — one page

YOUR AGENT, WIRED UP — ONE PICTURE

Every piece, and what it's for

YOU WRITE THESE ONCE — THEY BARELY CHANGE



TOGETHER — CHECKPOINT 1

Open the project. **Meet your week.**

- 01 **Copy the starter project** to your laptop, then open the folder in Claude Code. Just ask it: *"clone `github.com/nenti/weekly-briefing-agent` into a new folder."* One line, we do it together.
- 02 **Look at your week.** Open `data/inbox.md` and `data/calendar.md`. This is the sample inbox and calendar your agent will work through.

TOGETHER — CHECKPOINT 2

Now **build your agent**

03 **Drop in your About me.** Paste the profile Claude wrote for you at home into `me.md`.

04 **Build it.** Paste the build prompt. It asks you six questions — what your agent should do, how strict it should be, whether it may ever send anything — and writes `CLAUDE.md` from your answers. **That file is your agent.**

Nobody handed you that file — **you wrote it, by answering questions.** That is the whole skill.

TOGETHER — CHECKPOINT 3

Run it. Then turn it into **a page you'd show someone.**

05 Set your rules. `rubric.md` — what must always reach you, what's safe to skip. Plain sentences, no code.


06 Run it. Ask: *"Read my inbox and calendar, apply my rubric, build my weekly briefing."* Everyone now has a working agent.

07 Make it beautiful — and tune it. Ask for a polished HTML page, then open `briefing.html`. Not happy with a call it made? Tweak your rubric, run again. **That's iterating.**

If you only reach checkpoint 6, you still built a working agent tonight. **Nobody leaves empty-handed.**

THE PAYOFF

This is what you built



In two hours, your agent now:

- ✓ Reads your week — **email and calendar**
- ✓ Sorts what needs you from the noise, **by your rules**
- ✓ Drafts the replies that matter, **ready to send**
- ✓ Flags the decisions and the **calendar clash**
- ✓ Will do it again next Monday, **on its own**

No code. You told it what mattered — **it did the rest.**

THREE UPGRADES — IN THIS ORDER

Take it off the sample data

A • Your real week

Connect the Gmail + Calendar plug (one click, in your browser). Change one line — sample files to **your last 7 days**. Re-run. Use the throwaway account first.

B • Every Monday

`/schedule` a run on your own machine — the briefing is waiting before you sit down.


C • Laptop shut

Push it to GitHub and make it a **cloud routine** — it runs on Anthropic's infrastructure while your machine is off.

Start at A and work down. **Each one earns the next** — don't hand an agent the keys before you've watched it work.

A PRODUCTION AGENT — SAME THREE MOVES, A YEAR OF ITERATION

One line in. A whole **customer package** out.



Same **read** → **judge** → **write**, pointed at a bigger job. I'll show you a real one running — **then we build a small version of it together.**

TOGETHER — IF YOU'RE THROUGH THE BRIEFING

Build one that **talks**

- 01 **Pick any company** — yours, a client's, one you admire. Give the agent the name and the website.
- 02 **Let it do the reading.** It looks at the site, works out what they do and who they sell to, and writes a short script.
- 03 **Give it a voice.** ElevenLabs turns the script into a real voiceover — that's a tool with an **API key**, exactly like we just talked about.
- 04 **Out comes a video** you can actually send someone.

Different job, **same three moves**. Once you see that, you can point an agent at anything.

SAME THREE MOVES — AIMED AT ANOTHER JOB

What else you could **build**



Books your meetings

schedules

Say "set up 30 minutes with Tobi next week." It knows your contacts, finds a free slot in both calendars, and sends the invite — done.



Auto-drafted replies

drafts

When mail lands, it pulls context from your notes and past threads and writes the **full** answer, not just a stub.



Hands-off replies

sends

For safe, well-understood messages it answers **on its own** — inside the guardrails you set.



CRM in sync

syncs

An inbound signal lands in your **CRM** automatically, so your whole team sees it too.



Invoice to paid

acts

It reads the invoice, files it in your **ERP**, and gets it paid.

Every one of these is tonight's **read → judge → write**, aimed somewhere new.

THE LEAP — FROM YOU, TO YOUR WHOLE TEAM

Local is **your** agent. Remote moves your **company**.



Local · your personal agent

Runs on your laptop, in the loop, just for you. You watch every run and fix it in seconds. Perfect to build — and to learn to trust it.



Remote · your company's agent

Runs on a server for your **whole team**, on a trigger, without anyone watching. This is where an agent starts to move the business.

Honest: a cloud routine really does run with your laptop shut — but it only sees the connectors you set up at **claude.ai**, never what you connected locally. Closing that gap is the actual work.

WHERE PERSONAL BECOMES PRODUCTION

You deploy — agents run your software

nyxory

One platform you deploy your application to. From there the system takes care of operating your software — **with everything you need**.

nyx[®]argus

The back-office that **watches every run and scores it** — so you can trust the thousand runs nobody is looking at.

Monitoring

Backups

Incidents

Updates

24/7 operations

Same pattern as tonight, one level up: your coding agent hands the finished software to an **operations agent** — agent to agent, no ticket in between.

FOR YOU — AND A FRIEND

3 months of **nyxory**, on us



scan — the code comes pre-applied

Deploy the thing your agent builds. The **Starter plan** is **free for 3 months** — worth 147 euros.

`console.nyxory.com` • code **COWORK**

Send the link to a friend — it works for them too.

Scan, sign up, pick **Starter** — the discount attaches at checkout. Cancel any time.

YOU DID THIS TONIGHT

You built an agent.



Tobias Nentwig

Founder, nyxory — agents that operate, not just answer

linkedin.com/in/nentwig



let's connect

EVERYTHING FROM TONIGHT, IN ONE PLACE

It's all in **the repo**

github.com/nenti/weekly-briefing-agent

Do it again

workshop/guide.html

the whole workshop as a self-paced walkthrough — open it in a browser, redo any step at home

Go further

workshop/start-your-own-agent.md

for your NEXT agent — your own idea, empty folder to running agent, no starter repo

Prompt

workshop/about-me-prompt.md

the interview that writes your profile — useful in any AI tool, not just tonight

Prompt

workshop/build-your-agent-prompt.md

the interview from tonight — writes THIS agent's IDENTITY.md from your answers

Safety net

workshop/IDENTITY.example.md

a finished identity — compare, or copy it if you got stuck

The target

example/briefing-example.html

what a good briefing looks like — the standard your agent aims at

If you open one thing: **guide.html**. It is this whole evening, at your own pace.